Learning mdps from features: Predict-then-optimize for sequential decision making by reinforcemen...

Search within citing articles

INDEPENDENT

Decision-focused learning: Foundations, state of the art, benchmark and future opportunities

J Mandi, J Kotary, S Berden, M Mulamba… - Journal of Artificial …, 2024 - jair.org

Decision-focused learning (DFL) is an emerging paradigm that integrates machine learning (ML) and constrained optimization to enhance decision quality by training ML models in an …

Save Cite Cited by 222 Related articles All 12 versions

INDEPENDENT

Predict-and-optimize techniques for data-driven optimization problems: A review

A Anis Lahoud, AS Khan, E Schaffernicht… - Neural Processing …, 2025 - Springer

Abstract Machine learning predictive models rely on data to make predictions for new input data. However, accurate predictions are not always the end goal; practitioners often aim to …

Save Cite Cited by 5 Related articles All 3 versions

INDEPENDENT

First-order methods for linearly constrained bilevel optimization

G Kornowski, S Padmanabhan… - Advances in neural …, 2024 - proceedings.neurips.cc

Algorithms for bilevel optimization often encounter Hessian computations, which are prohibitive in high dimensions. While recent works offer first-order methods for …

Save Cite Cited by 20 Related articles All 8 versions

INDEPENDENT

A predict-then-optimize couriers allocation framework for emergency last-mile logistics

K Xia, L Lin, S Wang, H Wang, D Zhang… - Proceedings of the 29th …, 2023 - dl.acm.org

In recent years, emergency last-mile logistics (ELML) have played an essential role in urban emergencies. The efficient allocation of couriers in ELML is of practical significance to …

Save Cite Cited by 35 Related articles All 3 versions

Leaving the nest: Going beyond local loss functions for predict-then-optimize

S Shah, B Wilder, A Perrault, M Tambe - Proceedings of the AAAI …, 2024 - ojs.aaai.org

Predict-then-Optimize is a framework for using machine learning to perform decision-making under uncertainty. The central research question it asks is," How can we use the structure of …

Save Cite Cited by 32 Related articles All 5 versions

INDEPENDENT

Integrating symbolic reasoning into neural generative models for design generation

MJ Jacobson, Y Xue - Artificial Intelligence, 2025 - Elsevier

Abstract Design generation requires tight integration of neural and symbolic reasoning, as good design must meet explicit user needs and honor implicit rules for aesthetics, utility, and …

Save Cite Cited by 16 Related articles All 9 versions

Scalable decision-focused learning in restless multi-armed bandits with application to maternal and child health

K Wang, S Verma, A Mate, S Shah, A Taneja… - Proceedings of the …, 2023 - ojs.aaai.org

This paper studies restless multi-armed bandit (RMAB) problems with unknown arm transition dynamics but with known correlated arm features. The goal is to learn a model to …

Save Cite Cited by 42 Related articles All 8 versions

Restless multi-armed bandits for maternal and child health: Results from decision-focused learning

S Verma, A Mate, K Wang, N Madhiwalla… - Proceedings of the …, 2023 - ifaamas.org

ABSTRACT Mobile Health Awareness programs in underserved communities often suffer from diminishing engagement over time and health workers have to make live service calls …

Save Cite Cited by 27 Related articles All 4 versions

INDEPENDENT

A unified view on solving objective mismatch in model-based reinforcement learning

R Wei, N Lambert, A McDonald, A Garcia… - arXiv preprint arXiv …, 2023 - arxiv.org

Model-based Reinforcement Learning (MBRL) aims to make agents more sample-efficient, adaptive, and explainable by learning an explicit model of the environment. While the …

Save Cite Cited by 15 Related articles All 3 versions

INDEPENDENT

Sequential dynamic resource allocation in multi-beam satellite systems: A learning-based optimization method

Y Huang, WU Shufan, Z Zhankui, K Zeyu… - Chinese Journal of …, 2023 - Elsevier

Multi-beam antenna and beam hopping technologies are an effective solution for scarce satellite frequency resources. One of the primary challenges accompanying with Multi-Beam …

Save Cite Cited by 22 Related articles All 2 versions

Learning mdps from features: Predict-then-optimize for sequential decision making by reinforcemen...

Search within citing articles

Group fairness in predict-then-optimize settings for restless bandits

S Verma, Y Zhao, S Shah, N Boehmer… - The 40th Conference …, 2024 - openreview.net

Restless multi-arm bandits (RMABs) are a model for sequentially allocating a limited number of resources to agents modeled as Markov Decision Processes. RMABs have applications in …

Save Cite Cited by 10 Related articles All 5 versions

INDEPENDENT

Expert-calibrated learning for online optimization with switching costs

P Li, J Yang, S Ren - Proceedings of the ACM on Measurement and …, 2022 - dl.acm.org

We study online convex optimization with switching costs, a practically important but also extremely challenging problem due to the lack of complete offline information. By tapping …

Save Cite Cited by 21 Related articles All 7 versions

INDEPENDENT

Near-optimality for infinite-horizon restless bandits with many arms

X Zhang, PI Frazier - arXiv preprint arXiv:2203.15853, 2022 - arxiv.org

Restless bandits are an important class of problems with applications in recommender systems, active learning, revenue management and other areas. We consider infinite …

Save Cite Cited by 24 Related articles All 2 versions

INDEPENDENT

Decaf: Learning to be fair in multi-agent resource allocation

A Kumar, W Yeoh - arXiv preprint arXiv:2502.04281, 2025 - arxiv.org

A wide variety of resource allocation problems operate under resource constraints that are managed by a central arbitrator, with agents who evaluate and communicate preferences …

Save Cite Cited by 7 Related articles All 7 versions

INDEPENDENT

On solution functions of optimization: Universal approximation and covering number bounds

M Jin, V Khattar, H Kaushik, B Sel, R Jia - Proceedings of the AAAI …, 2023 - ojs.aaai.org

We study the expressibility and learnability of solution functions of convex optimization and their multi-layer architectural extension. The main results are:(1) the class of solution …

Save Cite Cited by 14 Related articles All 4 versions

INDEPENDENT

Conformal robust control of linear systems

Y Patel, S Rayan, A Tewari - arXiv preprint arXiv:2405.16250, 2024 - arxiv.org

End-to-end engineering design pipelines, in which designs are evaluated using concurrently defined optimal controllers, are becoming increasingly common in practice. To …

Save Cite Cited by 8 Related articles All 4 versions

INDEPENDENT

Learning for edge-weighted online bipartite matching with robustness guarantees

P Li, J Yang, S Ren - International Conference on Machine …, 2023 - proceedings.mlr.press

Many problems, such as online ad display, can be formulated as online bipartite matching. The crucial challenge lies in the nature of sequentially-revealed online item information …

Save Cite Cited by 14 Related articles All 8 versions

INDEPENDENT

Deep Reinforcement Learning in the Era of Foundation Models: A Survey

ID Mienye, E Esenogho, C Modisane - Computers, 2026 - mdpi.com

Deep reinforcement learning (DRL) and large foundation models (FMs) have reshaped modern artificial intelligence (AI) by enabling systems that learn from interaction while …

Save Cite Cited by 2 Related articles All 3 versions

INDEPENDENT

Utility-directed conformal prediction: A decision-aware framework for actionable uncertainty quantification

S Cortes-Gomez, C Patino, Y Byun, S Wu… - arXiv preprint arXiv …, 2024 - arxiv.org

Interest has been growing in decision-focused machine learning methods which train models to account for how their predictions are used in downstream optimization problems …

☆ Save 99 Cite Cited by 7 Related articles All 7 versions »

INDEPENDENT

Robustified learning for online optimization with memory costs

P Li, J Yang, S Ren - IEEE INFOCOM 2023-IEEE Conference …, 2023 - ieeexplore.ieee.org

Online optimization with memory costs has many real-world applications, where sequential actions are made without knowing the future input. Nonetheless, the memory cost couples …

☆ Save 99 Cite Cited by 10 Related articles All 4 versions

Learning mdps from features: Predict-then-optimize for sequential decision making by reinforcemen...

Search within citing articles

INDEPENDENT

Transfer reinforcement learning for mixed observability markov decision processes with time-varying interval-valued parameters and its application in pandemic …

M Du, H Yu, N Kong - INFORMS Journal on Computing, 2025 - pubsonline.informs.org

We investigate a novel type of online sequential decision problem under uncertainty, namely mixed observability Markov decision process with time-varying interval-valued parameters …

Save Cite Cited by 3 Related articles All 5 versions

INDEPENDENT

Decision-focused evaluation of worst-case distribution shift

K Ren, Y Byun, B Wilder - arXiv preprint arXiv:2407.03557, 2024 - arxiv.org

Distribution shift is a key challenge for predictive models in practice, creating the need to identify potentially harmful shifts in advance of deployment. Existing work typically defines …

Save Cite Cited by 6 Related articles All 6 versions

INDEPENDENT

A Learning-Based Energy-Efficient Device Grouping Mechanism for Massive Machine-Type Communication in the Context of Beyond 5G Networks

R Boisguene, I Althamary, CW Huang - Journal of Sensor and Actuator …, 2024 - mdpi.com

With the increasing demand for high data rates, low delay, and extended battery life, managing massive machine-type communication (mMTC) in the beyond 5G (B5G) context is …

Save Cite Cited by 3 Related articles All 3 versions

INDEPENDENT

A Review of End-to-End Decision Optimization Research: An Architectural Perspective

W Zhang, G Li - Algorithms, 2026 - mdpi.com

Traditional decision optimization methods primarily focus on model construction and solution, leaving parameter estimation and inter-variable relationships to statistical research …

Save Cite Related articles All 2 versions

INDEPENDENT

Online learning with real-time improving predictions

M Wu, Y Yin, JP Lynch - Artificial Intelligence for Transportation, 2026 - Elsevier

In this paper, we introduce the Decoupling Optimistic Online Mirror Descent (DOOMD) algorithm, a novel online learning approach designed for episodic Markov Decision …

Save Cite Related articles

INDEPENDENT

Learning-Augmented Online Decision Making: Algorithms, Analysis, and Applications

P Li - ACM SIGMETRICS Performance Evaluation Review, 2026 - dl.acm.org

Modern sequential decision-making systems, such as datacenter demand response and online ad allocation, operate under significant uncertainty while requiring strict adherence to …

Save Cite Related articles All 4 versions

INDEPENDENT

Policy-Driven World Model Adaptation for Robust Offline Model-based Reinforcement Learning

J Chen, L Xu, A Venugopal, J Schneider - arXiv preprint arXiv:2505.13709, 2025 - arxiv.org

Offline reinforcement learning (RL) offers a powerful paradigm for data-driven control. Compared to model-free approaches, offline model-based RL (MBRL) explicitly learns a …

Save Cite Related articles All 2 versions

INDEPENDENT

Differentiable arbitrating in zero-sum markov games

J Wang, M Song, F Gao, B Liu, Z Wang… - arXiv preprint arXiv …, 2023 - arxiv.org

We initiate the study of how to perturb the reward in a zero-sum Markov game with two players to induce a desirable Nash equilibrium, namely arbitrating. Such a problem admits a …

Save Cite Cited by 6 Related articles All 6 versions

INDEPENDENT

Applying reinforcement learning in Bitcoin trading to select technical strategies based on Deep Q-Network

NTT Hoan, NN Khang, PV Khanh - Cogent Economics & Finance, 2025 - Taylor & Francis

This paper proposes a Deep Q-Network-based reinforcement learning framework to optimize the selection of technical trading strategies for Bitcoin. Instead of directly predicting …

Save Cite Related articles All 4 versions

Case study: Applying decision focused learning in the real world

S Verma, A Mate, K Wang, A Taneja… - … on Trustworthy and …, 2022 - openreview.net

Many real world optimization problems with underlying unknown model parameters are solved using the predict-then-optimize framework. In particular, a model is learnt to first …

Save Cite Cited by 6 Related articles

Learning mdps from features: Predict-then-optimize for sequential decision making by reinforcemen...

Search within citing articles

INDEPENDENT

Timing is Important: Risk-aware Fund Allocation based on Time-Series Forecasting

F Lyu, L Du, Y Weng, Q Ying, Z Xu, W Zou… - Proceedings of the 31st …, 2025 - dl.acm.org

Fund allocation has been an increasingly important problem in the financial domain. In reality, we aim to allocate the funds to buy certain assets within a certain future period. Naive …

Save Cite Related articles All 5 versions

INDEPENDENT

MDP modeling for multi-stage stochastic programs

DP Morton, O Dowson, BK Pagnoncelli - arXiv preprint arXiv:2509.22981, 2025 - arxiv.org

We study a class of multi-stage stochastic programs, which incorporate modeling features from Markov decision processes (MDPs). This class includes structured MDPs with …

Save Cite Related articles All 4 versions

INDEPENDENT

Decisions and Deployment: The Five-Year SAHELI Project (2020-2025) on Restless Multi-Armed Bandits for Improving Maternal and Child Health

S Verma, A Dasgupta, N Madhiwalla, A Taneja… - arXiv preprint arXiv …, 2026 - arxiv.org

Maternal and child health is a critical concern around the world. In many global health programs disseminating preventive care and health information, limited healthcare worker …

Save Cite Related articles

INDEPENDENT

Cycles optimization to underwater glider standoff tracking profiles with the objective of conserving energy

W Zang, X Luan, D Song - Advanced Engineering Informatics, 2023 - Elsevier

Optimizing the energy consumption of electric powered underwater glider is critical. The objective of this work is to design the energy optimal cycles for standoff tracking profiles …

Save Cite Cited by 4 Related articles All 2 versions

INDEPENDENT

Decision-Focused Model-based Reinforcement Learning for Reward Transfer

A Sharma, S Parbhoo, O Gottesman… - arXiv preprint arXiv …, 2023 - arxiv.org

Decision-focused (DF) model-based reinforcement learning has recently been introduced as a powerful algorithm that can focus on learning the MDP dynamics that are most relevant for …

Save Cite Cited by 3 Related articles All 6 versions

Neural Index Policies for Restless Multi-Action Bandits with Heterogeneous Budgets

HS Pandey, K Wang, GGP Garcia - arXiv preprint arXiv:2510.22069, 2025 - arxiv.org

Restless multi-armed bandits (RMABs) provide a scalable framework for sequential decision-making under uncertainty, but classical formulations assume binary actions and a single …

Save Cite Related articles All 2 versions

INDEPENDENT

Using Operational Data Analytics for Planning Decisions Under Uncertainty

N Gautam - INFORMS Journal on Data Science, 2025 - pubsonline.informs.org

We consider a planning problem where events occur according to a renewal process with unknown parameters for the interevent times. We wish to make a decision to prepare for the …

Save Cite Related articles

INDEPENDENT

Simpo: Simultaneous prediction and optimization

B Zhang, YJ Ong, T Nakamura - 2022 IEEE International …, 2022 - ieeexplore.ieee.org

Many machine learning (ML) models are integrated within the context of a larger system as part of a key component for decision-making processes. Concretely, predictive ML models …

Save Cite Cited by 6 Related articles All 5 versions

Efficient Public Health Intervention Planning Using Decomposition-Based Decision-focused Learning

S Shah, A Suggala, M Tambe, A Taneja - arXiv preprint arXiv:2403.05683, 2024 - arxiv.org

The declining participation of beneficiaries over time is a key concern in public health programs. A popular strategy for improving retention is to have health workersintervene'on …

☆ Save 99 Cite Cited by 1 Related articles All 5 versions »

INDEPENDENT

Visualizing the implicit model selection tradeoff

Z He, Y Shaposhnik - Journal of Artificial Intelligence Research, 2023 - jair.org

The recent rise of machine learning (ML) has been leveraged by practitioners and researchers to provide new solutions to an ever growing number of business problems. As …

☆ Save 99 Cite Cited by 6 Related articles All 5 versions »

Learning mdps from features: Predict-then-optimize for sequential decision making by reinforcemen...

Search within citing articles

INDEPENDENT

A Novel Intelligent Decision Support System for Acute Coronary Syndrome Classification Using Reinforcement Learning-Enhanced Random Forests

RQ Ji, E Sejdić, K Kraevsky-Phillips, NT Riek… - … on Intelligent and Fuzzy …, 2025 - Springer

This study explores the integration of reinforcement learning (RL) with decision tree models to enhance the classification of acute coronary syndrome (ACS) using 12-lead ECG data …

Save Cite Related articles

INDEPENDENT

Near-optimality for multi-action multi-resource restless bandits with many arms

X Zhang - 2022 - search.proquest.com

We consider multi-action restless bandits with multiple resource constraints, also referred to as weakly coupled Markov decision processes. This problem is important in recommender …

Save Cite Cited by 3 Related articles All 3 versions

INDEPENDENT

Breaking the Scalability Barrier in Constrained Graph-Based Networked Control via Decision-Focused Learning

Z Yang, Y Guo, W Li, G Fan, H Jin, L Kong - Proceedings of the ACM …, 2026 - dl.acm.org

Many real-world systems can be modeled as graphs, where nodes store and consume entities, actively produce them, or have them emerge naturally, and edges transport them …

Save Cite Related articles

Decision-Focused Learning for the Masses With Applications to Public Health

S Shah - 2025 - search.proquest.com

In this thesis, I aim to better understand how to best train Machine Learning models for decision-making under uncertainty. Specifically, I focus on the" Predict-then …

Save Cite Related articles All 4 versions

INDEPENDENT

Task-Relevant Generative Models and Safe Reinforcement Learning with Applications to Clinical Decision-Making

A Sharma - 2025 - search.proquest.com

Clinical decision-making is inherently complex and high-stakes, frequently requiring sequential decisions under uncertainty. Automated discovery of clinical concepts and …

Save Cite Related articles All 2 versions

INDEPENDENT

Towards Fair Sequential Resource Allocation: Algorithmic Designs, Interventions, and Evaluations

A Kumar - 2025 - search.proquest.com

This thesis develops a comprehensive framework for fair sequential resource allocation in multi-agent systems where a centralized allocator coordinates actions under global …

Save Cite Related articles All 2 versions

INDEPENDENT

The Role of Explainable Artificial Intelligence in Model Selection, Adoption, and Wait Time Communication

Z He - 2025 - search.proquest.com

In this dissertation, I explore how people make decisions with AI, focusing on the role of explainable and interpretable AI in shaping these decisions. Across three chapters, I tackle …

Save Cite Related articles

INDEPENDENT

Overcoming data and model limitations in practical traffic prediction for intelligent transportation systems

M Li - 2025 - theses.lib.polyu.edu.hk

Intelligent Transportation Systems (ITS) leverage advanced technologies to enhance the efficiency, safety, and sustainability of transportation networks. A critical component of ITS is …

Save Cite Related articles All 2 versions

INDEPENDENT

Optimization in Deep Learning: Loss Landscape, Optimizer Dynamics, and Bi-Level Settings

J Wang - 2024 - search.proquest.com

Abstract Machine learning has witnessed remarkable advancements in recent years, transforming various industries and domains. Central to the success of machine learning …

Save Cite Related articles

INDEPENDENT

Health Impact of AI

Z Wu - 2025 - search.proquest.com

AI's explosive growth has triggered massive data center expansion with largely unexamined public health consequences. While carbon emissions receive attention, this paper quantifies …

Save Cite Related articles All 3 versions

Learning mdps from features: Predict-then-optimize for sequential decision making by reinforcemen...

Search within citing articles

INDEPENDENT

Decision-Focused Evaluation: Analyzing Performance of Deployed Restless Multi-Arm Bandits

P Verma, S Verma, A Mate, A Taneja… - arXiv preprint arXiv …, 2023 - arxiv.org

Restless multi-arm bandits (RMABs) is a popular decision-theoretic framework that has been used to model real-world sequential decision making problems in public health, wildlife …

Save Cite Cited by 2 Related articles All 3 versions

INDEPENDENT

Learning-Augmented Online Decision Making With Guaranteed Trustworthiness

J Yang - 2023 - search.proquest.com

Many mission-critical systems need to solve online decision-making problems such as workload scheduling in datacenters, power allocation in edge computing, battery …

Save Cite Related articles All 2 versions

Integrating Machine Learning and Optimization with Applications in Public Health and Sustainability

K Wang - 2023 - search.proquest.com

The field of artificial intelligence (AI) has garnered increasing attention in the realms of public health and conservation due to its potential to characterize complex dynamics and …

Save Cite Related articles All 2 versions

INDEPENDENT

Inverse Reinforcement Learning for Restless Multi-Armed Bandits with Application to Maternal and Child Health

G Jain, P Varakantham, H Xu, A Taneja… - BridgeAICCHE2023 …, 2023 - openreview.net

We study restless multi-armed bandits (RMABs) in the context of public health, where there is a need to optimize resource allocation decisions. Until now, RMABs typically solve for the …

Save Cite Related articles All 2 versions

INDEPENDENT

Learning with Real-time Improving Predictions in Online MDPs

M Wu, Y Yin, J Lynch - openreview.net

In this paper, we introduce the Decoupling Optimistic Online Mirror Descent (DOOMD) algorithm, a novel online learning approach designed for episodic Markov Decision …

Save Cite Related articles

INDEPENDENT

Temporal analysis of healthcare trajectories using deep learning: a dissertation in Engineering and Applied Science

R Ryan - 2023 - repository.lib.umassd.edu

Abstract Electronic Health Record (EHR) has gained significant importance in recent years due to its potential aspects in clinical prognosis and treatment. Patients' clinical information …

Save Cite Related articles All 2 versions

INDEPENDENT

Decision-Focused Offline Deep Reinforcement Learning for Healthcare Policy Optimization

P Bhoj, A Eshragh, Y Li - cs224r.stanford.edu

Method. We couple a neural predictor ϕv of transition dynamics with a Double Deep-Q Network (DDQN) planner, forming a differentiable end-to-end pipeline. Policy values are …

Save Cite Related articles

Robust decision-focused learning for reward transfer

A Sharma, S Parbhoo, O Gottesman, F Doshi-Velez - arXiv preprint arXiv:2304.03365, 2023

Save Cite Cited by 1 Related articles